

```
[2]: #import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
import warnings
warnings.filterwarnings("ignore")

#Load and inspect dataset
df=pd.read_csv("sales_data_.csv",parse_dates=['Date'],index_col='Date')
df = df.asfreq('D')

#display first few rows
print(df.head())

#check for missing values
print("\nMissing Values:")
print(df.isnull().sum())

#visualize sales trends
#plot sales over time
plt.figure(figsize=(10,5))
plt.plot(df.index,df['Sales'],label="Sales Data",color='blue')
plt.xlabel('Date')
plt.ylabel('Sales')
plt.title('Sales Trend Over Time')
plt.legend()
plt.show()
```



```
#Check Stationarity (Dickey-Fuller Test)
from statsmodels.tsa.stattools import adfuller

#Perform Dickey-Fuller Test
result=adfuller(df['Sales'])
print(f"ADF Statistic: {result[0]}")
print(f"P-Value: {result[1]}")

if result[1]>0.05:
    print("The data is non-stationary. Differencing is needed.")
else:
    print("The data is stationary.")

#build ARIMA Model
#Fit ARIMA Model (assuming p=2,d=1,q=2 based on ACF/PACF)
model=ARIMA(df['Sales'],order=(2,1,2))
model_fit=model.fit()

#forecast next 10 periods
forecast=model_fit.forecast(steps=10)
print("Forecasted Sales for next 10 periods:\n",forecast)

#visualize forecast
#plot original and forecasted sales
plt.figure(figsize=(10,5))
plt.plot(df.index,df['Sales'],label="Actual Sales",color='blue')
plt.plot(pd.date_range(start=df.index[-1],periods=11,freq='D')[1:],forecast,label="Forecast",color='red')
plt.xlabel("Date")
plt.ylabel("Sales")
plt.title("Sales Forecast using ARIMA")
plt.legend()
plt.show()
```

```
#project 2
#importing necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

#Load and inspect dataset
df=pd.read_csv("heart_disease.csv")

#display first few rows
print(df.head())

#check for missing values
print("\nMissing Values:")
print(df.isnull().sum())

#preprocess data
#define features (X) and target variable (y)
X=df.drop(columns=['Heart Disease'])
y=df['Heart Disease']
X_encoded = pd.get_dummies(X, drop_first=True)

#standardize numerical features
scaler=StandardScaler()
X_scaled=scaler.fit_transform(X_encoded)

#split into train and test sets
```

```
#split into train and test sets
X_train,X_test,y_train,y_test=train_test_split(X_scaled,y,test_size=0.2,random_state=42)
print(f"Training samples: {X_train.shape[0]},Testing samples:{X_test.shape[0]}")

#train logistic regression model
model=LogisticRegression()
model.fit(X_train,y_train)

#predict on test data
y_pred=model.predict(X_test)

#print accuracy
accuracy=accuracy_score(y_test,y_pred)
print(f"ModelAccuracy:{accuracy:.2f}")

#evaluate model performance
#confusion matrix
cm=confusion_matrix(y_test,y_pred)
print("\nConfusion Matrix:")
print(cm)

#classification report
print("\nClassification Report:")
print(classification_report(y_test,y_pred))


#plot confusion matrix
plt.figure(figsize=(6,4))
sns.heatmap(cm,annot=True,fmt="d",cmap="Blues",xticklabels=["No Disease","Disease"],yticklabels=["No Disease","Disease"])
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```

| Date | Sales | Product | Region |
|------------|-------|---------|--------|
| 2024-01-01 | 150 | A | North |
| 2024-01-02 | 200 | B | South |
| 2024-01-03 | 180 | A | North |
| 2024-01-04 | 220 | C | East |
| 2024-01-05 | 160 | A | West |

Missing Values:
Sales 0
Product 0
Region 0
dtype: int64



Sales Trend Over Time



ADF Statistic: 1.896571585241519

P-Value: 0.9985224831326911


The data is non-stationary. Differencing is needed.

Forecasted Sales for next 10 periods:


| | |
|------------|------------|
| 2024-01-11 | 240.000047 |
| 2024-01-12 | 240.000038 |
| 2024-01-13 | 240.000038 |
| 2024-01-14 | 240.000038 |
| 2024-01-15 | 240.000038 |
| 2024-01-16 | 240.000038 |
| 2024-01-17 | 240.000038 |
| 2024-01-18 | 240.000038 |
| 2024-01-19 | 240.000038 |
| 2024-01-20 | 240.000038 |

Freq: D, Name: predicted_mean, dtype: float64

Sales Forecast using ARIMA



Sales Forecast using ARIMA



| | Age | Gender | Cholesterol | Blood Pressure | Heart Disease |
|---|-----|--------|-------------|----------------|---------------|
| 0 | 52 | Male | 245 | 130/85 | 1 |
| 1 | 58 | Female | 210 | 140/90 | 0 |
| 2 | 46 | Male | 190 | 120/80 | 0 |
| 3 | 59 | Female | 225 | 135/88 | 1 |
| 4 | 63 | Male | 250 | 150/95 | 1 |

Missing Values:

Age 0
Gender 0
Cholesterol 0
Blood Pressure 0
Heart Disease 0

dtype: int64

Training samples: 8, Testing samples: 2

Model Accuracy: 0.50

Confusion Matrix:

[[0 1]

[0 1]]

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.00 | 0.00 | 0.00 | 1 |
| 1 | 0.50 | 1.00 | 0.67 | 1 |
| accuracy | | | 0.50 | 2 |
| macro avg | 0.25 | 0.50 | 0.33 | 2 |
| weighted avg | 0.25 | 0.50 | 0.33 | 2 |

